

An Introduction to High Performance Computing

Stuart Rankin
sjr20@cam.ac.uk

Research Computing Services (<http://www.hpc.cam.ac.uk/>)
University Information Services (<http://www.uis.cam.ac.uk/>)

30th July 2019 / UIS Training

Welcome

- ▶ Please sign in on the attendance sheet.
- ▶ Keep your belongings with you.
- ▶ Please ask questions and let us know if you need assistance.

UIS: Research Computing Services

Your trainers for today will be:

- ▶ **Paul Sumption**
Research Computing User Services
- ▶ **Eleftherios Avramidis**
Research Software Engineering



**British
Antarctic Survey**
NATURAL ENVIRONMENT RESEARCH COUNCIL

**POLAR SCIENCE
FOR PLANET EARTH**

You may be ...

- ▶ Programmers (or not).
- ▶ UNIX power users (or not).
- ▶ Researchers wishing to run large, parallel code.
- ▶ Researchers wishing to run many, non-parallel cases.
- ▶ Researchers interested in big data, machine learning, AI.
- ▶ Researchers requiring slightly more than an ordinary workstation.
- ▶ Many different disciplines and requirements.

You may be ...

- ▶ Programmers (or not).
- ▶ UNIX power users (or not).
- ▶ Researchers wishing to run large, parallel code.
- ▶ Researchers wishing to run many, non-parallel cases.
- ▶ Researchers interested in big data, machine learning, AI.
- ▶ Researchers requiring slightly more than an ordinary workstation.
- ▶ Many different disciplines and requirements.

You may be ...

- ▶ Programmers (or not).
- ▶ UNIX power users (or not).
- ▶ Researchers wishing to run large, parallel code.
- ▶ Researchers wishing to run many, non-parallel cases.
- ▶ Researchers interested in big data, machine learning, AI.
- ▶ Researchers requiring slightly more than an ordinary workstation.
- ▶ Many different disciplines and requirements.

You may be ...

- ▶ Programmers (or not).
- ▶ UNIX power users (or not).
- ▶ Researchers wishing to run large, parallel code.
- ▶ Researchers wishing to run many, non-parallel cases.
- ▶ Researchers interested in big data, machine learning, AI.
- ▶ Researchers requiring slightly more than an ordinary workstation.
- ▶ Many different disciplines and requirements.

You may be ...

- ▶ Programmers (or not).
- ▶ UNIX power users (or not).
- ▶ Researchers wishing to run large, parallel code.
- ▶ Researchers wishing to run many, non-parallel cases.
- ▶ Researchers interested in big data, machine learning, AI.
- ▶ Researchers requiring slightly more than an ordinary workstation.
- ▶ Many different disciplines and requirements.

You may be ...

- ▶ Programmers (or not).
- ▶ UNIX power users (or not).
- ▶ Researchers wishing to run large, parallel code.
- ▶ Researchers wishing to run many, non-parallel cases.
- ▶ Researchers interested in big data, machine learning, AI.
- ▶ Researchers requiring slightly more than an ordinary workstation.
- ▶ Many different disciplines and requirements.

You may be ...

- ▶ Programmers (or not).
- ▶ UNIX power users (or not).
- ▶ Researchers wishing to run large, parallel code.
- ▶ Researchers wishing to run many, non-parallel cases.
- ▶ Researchers interested in big data, machine learning, AI.
- ▶ Researchers requiring slightly more than an ordinary workstation.
- ▶ Many different disciplines and requirements.

Plan of the Course

Part 1: Basics

Part 2: Research Computing Services HPC

Part 3: Using HPC

10:00 WELCOME

11:00-11:15 Break

12:30-13:30 LUNCH

15:30-15:45 Break

16:30 CLOSE

Plan of the Course

Part 1: Basics

Part 2: Research Computing Services HPC

Part 3: Using HPC

10:00 WELCOME

11:00-11:15 Break

12:30-13:30 LUNCH

15:30-15:45 Break

16:30 CLOSE

article amsmath color,pxfonts,fix-cm latexsym [mathletters]ucs
8211- 46. 58: 8226• 8217' 62> 32 [T1]fontenc [utf8x]inputenc
pict2e wasysym [english]babel tikz
[margin=0in,paperwidth=720pt,paperheight=540pt]geometry

Topics Covered

- Hardware

- Storage

- Access

- Data Transfer

Hardware

- Gateway or Bastion hosts (bslceb & bslceb)

- Only use for access to BAS or transferring files

- Headnodes

- No access, manages job queues and storage

- General Use Workstations & Private Workstations

A c c e s s

Authentication

- Three passwords • UNIX (NIS), LDAP and Samba

- UNIX for bslnb /bslnc and LDAP for HP

- Try to keep all these password synchronised

- We are working to simplify the situation

SSH

Access (continued)

X2Go

- Access HPC desktop interface with or without a mouse
- Disconnecting and reconnecting
- Copy/paste
- Sharing files from your laptop or PC
- More information: <http://ictdocs/wiki/index.php/X2Go>

Storage

User Area - /users/user name

- Small, not intended for sharing data

- Space restricted via quotas

SAN Volumes

- Setup for projects and departments, e.g.: /data

- Accessible from workstations, balances, balances

Storage (continued)

HPC Storage - /data/hpcdata/ (users, data)

- â€ Accessible from nodes and workstations, bsl

- â€ Usage limited via quotas

Quotas

- â€ On HPC you can check your quotas using: m

- â€ Need more space - contact the service desk

Data Access and Transfer

Samba

- Allows clients to connect to UNIX storage as if it was a Windows share
- Allows access to SAN volumes, /users and /data
- No access to /data/hpcflash

FTP

- Allows non-BAS users to retrieve files from the

Data Access and Transfer(c

â€¢rsync

â€¢ Perfect tool for transferring file locally and s

â€¢ Options to resume , reconnect, compression

â€¢scp

â€¢sshfs

User Environment

• Shell

- Our default shell is tcsh

- If you prefer something different such as bash

• SSH keys

- Connect to BAS systems without typing pass

- ssh-keygen • Always create with a passphrase

User Environment

tmux

- Keeps long running command line sessions

- Allows disconnecting and reconnecting

- Multiple command line sessions and console

- More information: <http://ictdocs/wiki/index.php>

- Demonstration

Software

Operating System software

- Typical linux commands and some graphical
- These can be run from the command line and

Modules

- Do not work on bslnb or bslnb
- There are two module repositories: /package

Software (continued)

Modules

â€¢ Useful module commands:

`module avail`

`module load name/version`

`module unload name/version`

â€¢ Common mistakes

Jupyter Notebooks

• Jupyter notebooks running on workstation

• More information: http://ictdocs/wiki/index.php/Jupyter_Notebooks

Containers

Containers at BAS are still a work in progress

What are containers?

Podman

To be able to use it, you need to contact the service

Container images must be downloaded to each node

SLURM

• What is it?

• Simple Linux Utility for Resource Management

• Schedules jobs based on the resources they

• Different queues

• Short -

• Medium

SLURM

Job Types

- Batch • Standard job

- MPI

- When you require large amounts of memory or cpu cores

- MPI require infiniband connectivity for Messaging

- All nodes need to be in the same queue.

SLURM (continued)

• Job submission scripts • Simple:

SLURM (continued)

• Job monitoring

• `squeue -u <username>`

• `sacct -j <jid>`

• To see details on resources used by all running jobs:

• To see all your recent jobs: `sacct -u <username>`

• To check memory and cpu usage on a node:

SLURM (continued)

- Troubleshooting failed jobs

- Set an output file, often has useful information

- Common mistakes

- Forgetting to load modules

- Using storage which is not visible to the HPC

- Avoid using symlinks

Model Ensemble

• Aimed at solving complex HPC workload

Developer Environment

• Use git!

• Repeatable, reproducible & shareable

Best Practice

â€¢ User Policy

â€¢ Link:

â€¢ Do!

â€¢ Ask for help,

â€¢ Don't!

HELP!

Service Desk: services.k@bas.ac.uk

HPC User Guide: <http://ictdocs/wiki/index.php/HPC>

Service Desk Solutions

Yammer



Any Questions?

The Following Ar



Title

Title

Title

Title

Title

Title

Title

Title

Title

Title