



# An Introduction to High Performance Computing 2021

Paul Sumption  
pasump@bas.ac.uk

December 20, 2021



# Exercise 1: Login

- Log into your RCS training account.

*Hints:* Create a terminal window and use **ssh** to login to your cluster training account.

The remote host is **login.hpc.cam.ac.uk**. The user name is the same name as your MCS Desktop training account (i.e. **z4XY**).

N.B. If in doubt about the name of your training account, check the number of your station (see the label on the top of the box), then station **1XY** should correspond to account **z4XY**.



## Exercise 2: Simple command line operations

- (a) List your current directory (folder) using **ls -al**. Use **df -h** to see the various cluster filesystems, their sizes and their current total usages. You will be on a random login node – use **hostname** to confirm which one, and **w** to find out who else is using it. Use **lstopo** to find out more about the internal structure of the login node.
- (b) Examine your personal filesystem quotas with the command **quota**.

You should see a 40GB quota on /home, a 1TB block and 1024k file quota on /rds-d2 (which corresponds to ~/rds/hpc-work).

- (c) Ask the scheduler what compute resources are available to you with **mybalance**. This command may take a little while to return (the units are CPU/GPU/KNL hours).



## Exercise 3: Laptop setup

Before attempting exercise 3:

- ▶ <https://www.hpc.cam.ac.uk/training-courses>
- ▶ Download the file **exercises.tgz** to your desktop.
- ▶ You will need an sftp client on your laptop to then transfer this file to the cluster.
- ▶ Mac users - open a terminal, you can use sftp from the command line
- ▶ Windows users - download winsftp



## Exercise 3: File transfer

- Use SFTP to transfer the file **exercises.tgz** to your Research Computing Service training account directory `~/rds/hpc-work`.

*Hints:* The command is **sftp**. Use the same remote host, username and password as in the previous exercise.

Use **cd rds/hpc-work** to change the target directory, then **put exercises.tgz** to transfer the file from your desktop to the target directory on the Research Computing Service cluster. Use **quit** to close the connection.

Optionally, copy the file over again using **rsync**.



## Exercise 3: File transfer (ctd)

- ▶ Switch back to the SSH session you created in the previous exercise. Verify that the file is now present by using **ls**.

*Hints:* Do **ls -al ~/rds/hpc-work/**. Note that you can often reduce typing by pressing **TAB**.

- ▶ Unpack the tar archive to create an exercise subdirectory.

*Hints:* Do **cd ~/rds/hpc-work/** then **tar -zxvf exercises.tgz**.



## Exercise 5: Modules and Compilers

- ▶ Go to the **exercises** directory of your cluster account.

*Hints:* Firstly you may need to review Exercise 1 in order to reconnect to your cluster account. At the remote command prompt, change to the exercises directory (`cd ~/rds/hpc-work/exercises`).

- ▶ Try to compile the **hello.c** program using the default **gcc** compiler (it will fail because there is a deliberate bug).

*Hints:* `gcc hello.c -o hello`

- ▶ To fix the problem, open the **hello.c** file in an editor (e.g. **gedit**, **nano**, **emacs**).

*Hints:* Launch gedit in the background by doing `gedit&`. A gedit window should appear. Remove the word **BUG**, save the file and recompile. Do `./hello` to run the program.



British  
Antarctic Survey

NATURAL ENVIRONMENT RESEARCH COUNCIL



POLAR SCIENCE  
FOR PLANET EARTH

## Exercise 5: Modules and Compilers (ctd)

- ▶ The default version of gcc on the RCS HPC clusters is 4.8.5. Compile hello.c again with **gcc 5.4.0**.

*Hints:* module av, module load, then **gcc hello.c -o hello2**

- ▶ Launch the Matlab GUI. Note this should work from either the SSH command-line or remote desktop sessions.

*Hints:* **module load matlab** then run: **matlab&**

- ▶ Quit Matlab and launch it again without the graphical desktop interface. This is the way to launch it inside a batch job.

*Hints:* **matlab -nodisplay -nojvm -nosplash**





## Exercise 6: Submitting Jobs (Matlab)

- Submit a job which will run **matlab** on the **file.m** command file (which contains just the Matlab **ver** command).

*Hints:*

1. Load the matlab module at the place indicated in the file **job\_script** in your exercises directory.
2. Set the value of application to **"matlab -nodesktop -nosplash -nojvm"**
3. Set the value of options to **"-r file"**
4. Submit the job with **sbatch job\_script**. The jobid is then printed.
5. Watch the job in the queue with **squeue**.
6. After it has disappeared, open the output file **slurm-jobid.out** in your editor. It should contain a list of licensed Matlab features from the ver command.
7. For more demanding work you can increase the available memory by increasing the number of cpus.



## Exercise 7: Submitting Jobs (serial or threaded application)

- ▶ Submit a job which will run a copy of your hello program on 1 cpu.

- Hints:*
1. Edit the script **job\_script** in your exercises directory. Set:  
**#SBATCH --nodes=1**  
**#SBATCH --ntasks=1**  
**application="./hello"**
  2. Submit the job with **sbatch job\_script**. The jobid is then printed.
  3. Watch the job in the queue with **squeue**.
  4. After it has disappeared, open the output file **slurm-jobid.out** in your editor. There should be exactly one "Hello, World!" message.

Experiment with varying the number of nodes and tasks (you are limited to 4 nodes). Note you will need to launch the application with **run** to actually use more than 1 cpu.



## Exercise 8: Submitting Jobs (R)

- ▶ R jobs may be serial, threaded, or even MPI parallel depending on the packages used. Submit a job which will run the trivial script **hello.r** program on 1 cpu.

*Hints:* 1. Edit the script **job\_script** in your exercises directory. Set:

```
#SBATCH --nodes=1  
#SBATCH --ntasks=1  
application="Rscript"  
options="hello.r"
```

2. Submit the job with **sbatch job\_script**. The jobid is then printed.

- ▶ Repeat this using a different version of R.



## Exercise 9: Array Jobs

- ▶ Submit your last job in the form of an array with indices 1-64. Use `-H` with `sbatch` to mark the array as held (so that it won't run immediately).

*Hints:*

1. Use **`sbatch -H --array=1-64 job_script`**
2. Use **`squeue -u userid`** to see your array job.  
Note that `-r` reports each array element individually.

- ▶ Release array element 1 and allow it to run. Then release the others.

*Hints:*

1. Use **`scontrol release`**  
**`${SLURM_ARRAY_JOB_ID}_1`**
2. Use **`squeue -u userid`** again to watch what happens.
3. Release the others with  
`scontrol release`  
**`${SLURM_ARRAY_JOB_ID}`**  
i.e. use the array id to release the entire array.
4. When all the jobs complete you should have 64  
`slurm-${SLURM_ARRAY_JOB_ID}_1` and 63  
saying hello from various cpus on possibly

